

Lec = 01 :- Complete theory of Kubernetes.

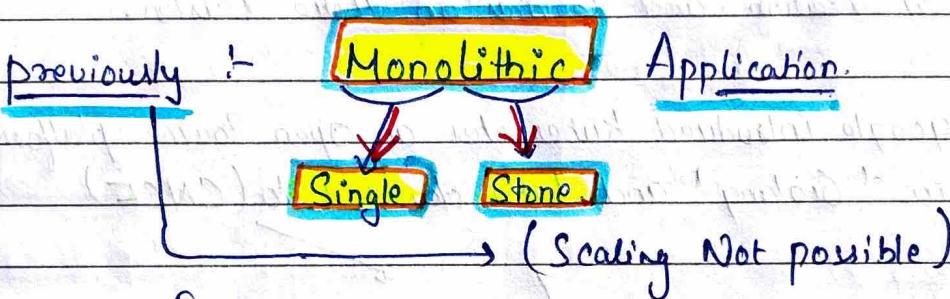


↓
(K8s)

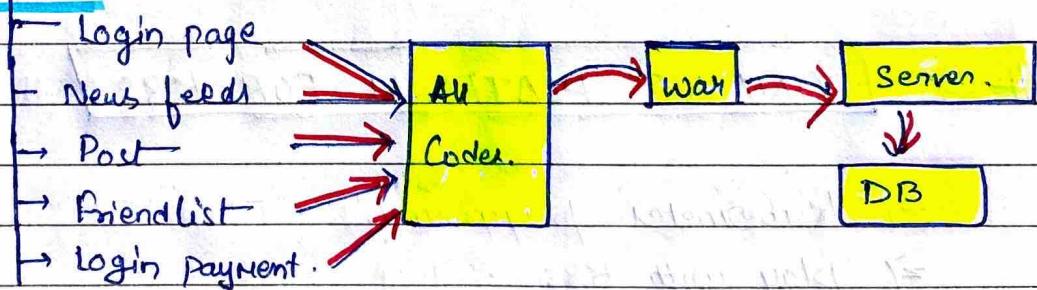
7 lines in the logo.
(project 7).

- ↙ (Container Management tool)
- ↳ (High Availability.)
- ↳ (Scalability)

Kubernetes is an open-source Container - Orchestration System for automating Computer application deployment, Scaling, and Management.



Facebook,

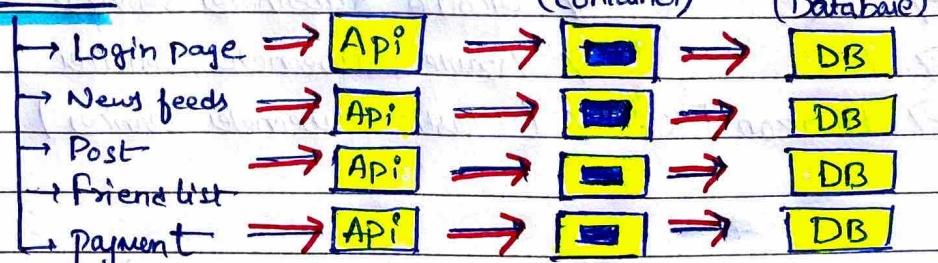


Now :-

Micro Services.

1 → (works independently.)

Facebook



⇒ Kubernetes is an Open-Source Container Management tool which automates Container deployment, Container scaling & Load Balancing.

⇒ It schedules, runs and manages isolated containers which are running on virtual / physical / cloud machines.

⇒ All top cloud providers support.

History

⇒ Google developed an internal system called 'borg' (later named as omega) to deploy and manage thousand google applications and services on their clusters.

⇒ In 2014, Google introduced Kubernetes an open source platform written in 'GoLang' and later donated to (CNCF)

(Cloud Native Computing Foundation)

⇒ ONLINE PLATFORM FOR K8S

⇒ Kubernetes playground.

⇒ play with K8s

⇒ play with kubennets classroom.

⇒ CLOUD BASED K8S SERVICES

⇒ GKE → Google Kubernetes Service.

⇒ AKS → Azure Kubernetes Service

⇒ Amazon EKS → (AWS Kubernetes Service)

7 KUBERNETS INSTALLATION TOOL

- ⇒ Minicube.
- ⇒ Kubeadm.

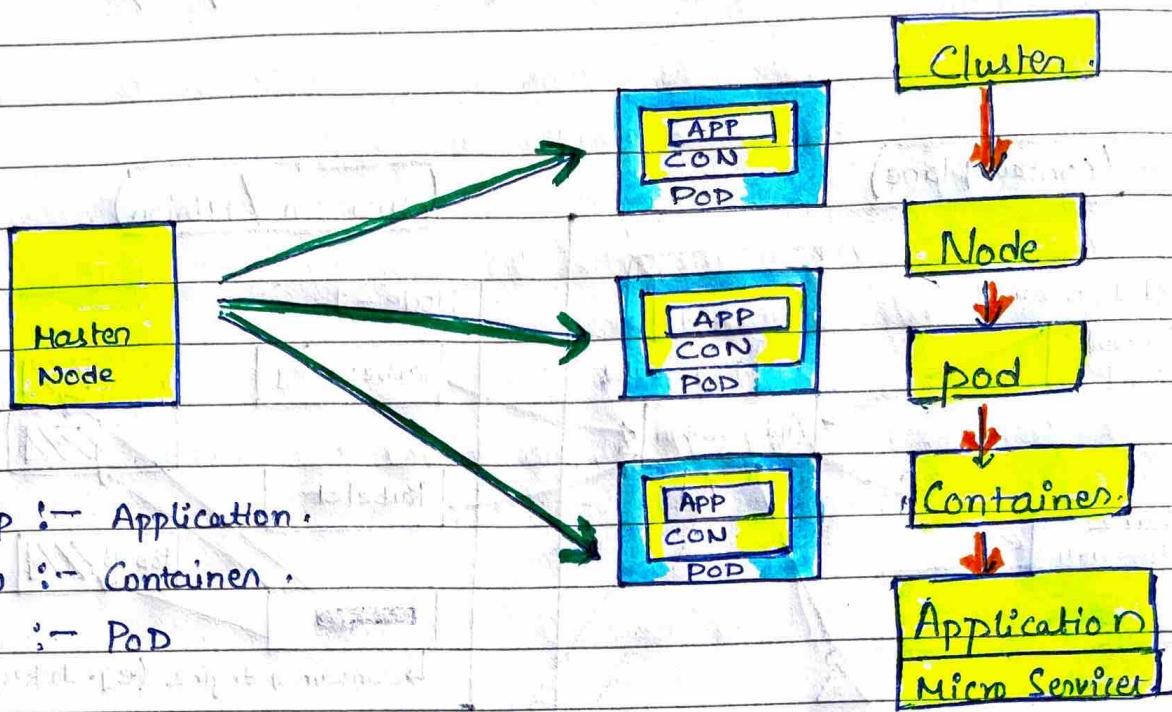
Problems with Scaling up the Containers.

- ⇒ Containers Can not communicate with each other.
- ⇒ Auto-scaling and Load Balancing was not possible.
- ⇒ Containers had to be managed carefully.

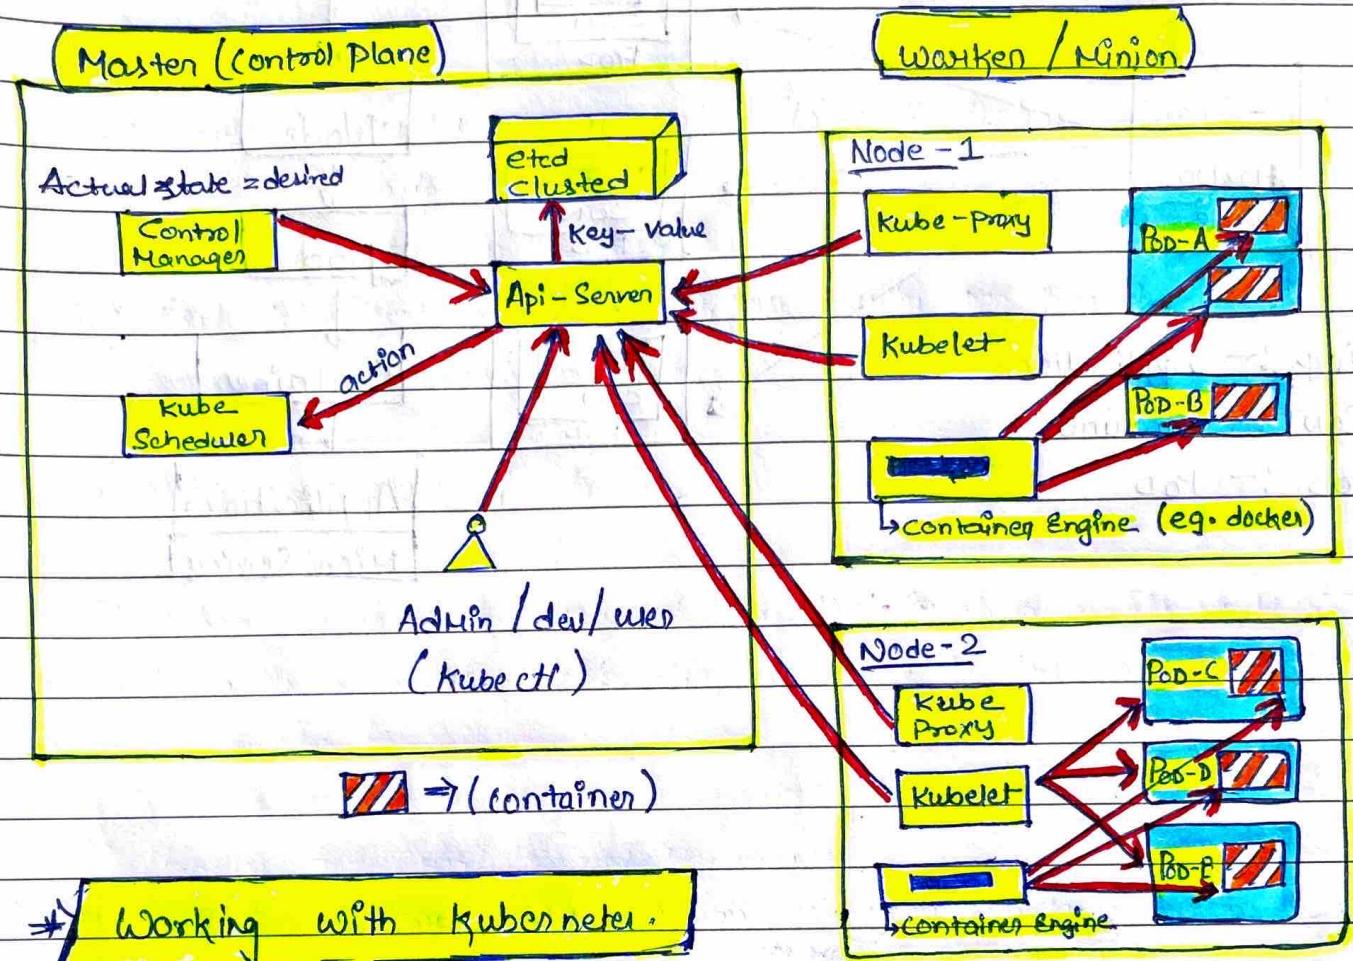
Features of Kubernetes.

- ⇒ Orchestration (clustering of any no of containers running on different v/w)
(vertical & horizontal) ↓ Prefend.
- ⇒ Auto Scaling ⇒ Auto - healing.
- ⇒ Load Balancing ⇒ platform independent (cloud / virtual / physical)
- ⇒ Fault Tolerance (Node / pod failure)
- ⇒ Rollback (going back to previous version)
- ⇒ Health Monitoring of Containers.
- ⇒ Batch Execution (one time Sequential, Parallel)

FEATURES	KUBERNETES	DOCKER SWARM
Installation and Cluster Configuration	Complicated and time consuming	Fast and Easy.
Supports	k8s Can work with almost all container types like Rocket, Docker Container	Work with docker only.
GUI	GUI Available	GUI Not Available.
Data Volumes	Only shared with containers in same pod.	Can be shared with any other container.
Update & Rollback	Process Scheduling to maintain services while updating	Progressive update & Service health monitoring throughout the update.
Auto Scaling	Support Vertical and Horizontal Auto scaling	Not support Auto scaling.
Logging and Monitoring	Inbuilt tool present for monitoring	Used 3rd party tools like Splunk.



Aec-02 Architecture of Kubernetes & its demo.



- ⇒ We Create Manifest (.yaml)
- ⇒ Apply this to Cluster (to Master) to bring into desired state
- ⇒ pod runs on node, which is controlled by Master.

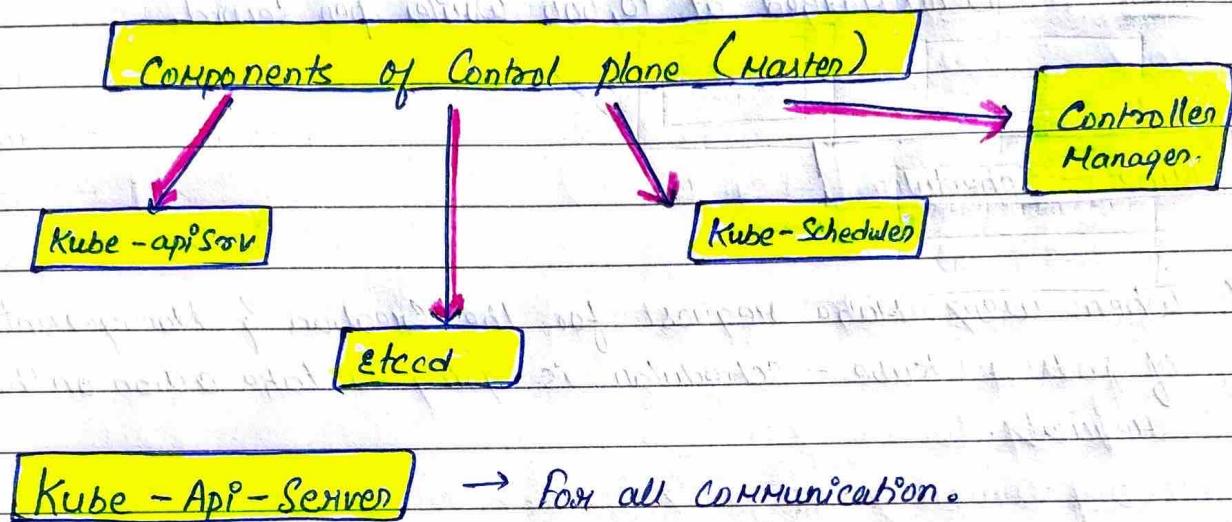
Role of Master Node.

- ⇒ Kubernetes cluster contains containers running on bare metal / VM instances / cloud instances / all mix.
- ⇒ Kubernetes designates one or more of these as Master and all others as Worker.

⇒ The Master is now going to run set of K8s processes. These processes will ensure smooth functioning of cluster. These processes are called "Control plane".

⇒ Can be Multi-Master for high availability.

⇒ Master runs Control plane to run cluster smoothly.



⇒ This API-server interacts directly with your user (i.e. we apply YAML or JSON manifest to Kube API-server).

⇒ This Kube-API server is meant to scale automatically as per load.

⇒ Kube API-server is front-end of Control-plane.

etcd

⇒ Stores metadata and status of cluster.

⇒ etcd is consistent and high availability store (key-value store).

⇒ Storehouse of truth for cluster state.

(information about the state of cluster).

etcd has following features :-

1. Fully Replicated :- The entire state is available on every node in the cluster.
2. Secure :- Implements automatic TLS with optional client - certificate authentication.
3. Fast :- Benchmarked at 10,000 writes per second.

Kube-Scheduler

- ⇒ When user make request for the creation & management of pods, Kube-Scheduler is going to take action on the requests.
- ⇒ Handles pod creation and management.
- ⇒ Kube-Scheduler match / assign any node Create and Run pods.
- ⇒ A scheduler watches for newly created pods that have no node assigned for every pod that the scheduler discovers, the scheduler becomes responsible for finding best node for that pod to run on.
- ⇒ Scheduler gets the information for hardware configuration .yaml file and schedules the pods on nodes accordingly.

Controller - Manager.

⇒ Make Sure actual state of Cluster Matcher to desired State.

Two possible choices for Controller Manager.

- ① if K8s On cloud, then it will be cloud - controller - Manager.
- ② if K8s ON Non-cloud then it will be Kube - controller - Manager.

Components on Master that runs Controller.

Node - Controller → For checking the Cloud provider to determine if a Node has been detected in the cloud after it's stop responding.

Route - Controller → Responsible for setting up network, routers on your cloud.

Service - Controller :— Responsible for load Balancers on your cloud against Services or types load Balancers.

Volume - Controller :— For Creating, attaching and Mounting volumes and interacting with the cloud provider to orchestrate volume.

Nodes (Kubelet and Container Engine)

⇒ Node is going to run 3 important piece of software / process.

Kubelet

- ⇒ Agent running on the node.
- ⇒ Listens to Kubernetes Master. (e.g. pod creation request)
- ⇒ use port **10255**
- ⇒ Sends Success and Fail report to Master.

Container Engine (e.g. Docker)

- ⇒ Works with Kubelet
- ⇒ pulling images.
- ⇒ Start / stop containers.
- ⇒ Exposing Containers on ports specified in Manifest.

Kube - Proxy

- ⇒ Assign IP to each pod.
- ⇒ It is required to assign (IP) address to pods (Dynamic IP)
- ⇒ Kube-proxy runs on each node & this make sure that each pod will get its own unique IP address.

These 3 Components Collectively Constitute "Node"

POD

- ⇒ Smallest unit in Kubernetes.
- ⇒ POD is a group of One or More Containers that are deployed together on the same host.
- ⇒ A cluster is a group of Nodes.
- ⇒ A cluster has at least one worker node and Master node.
- ⇒ In Kubernetes, the Control Unit is the pod, not Container.
- ⇒ Consist of one or more tightly coupled Containers.
- ⇒ POD runs on node, which is controlled by Master.
- ⇒ Kubernetes only knows about pods
(does not know about individual containers)
- ⇒ Cannot start Container without a pod.
- ⇒ One POD usually contains one controller.

Multicontainer Pods

- ⇒ Share access to memory space.
- ⇒ Connect to each other using (local host).
- ⇒ Share access to the same volume.
- ⇒ Containers within pods are - deployed in an all nothing manner.
- ⇒ Entire Pod is hosted on the same Node (Scheduler will decide about which Node).

POD LIMITATION

- ⇒ No auto-healing or scaling.
- ⇒ POD crashes.

Higher level k8s Object

Replication Set :- Scaling and auto healing.

Deployment :- Versioning and Rollback
Service :- Static IP and Networking
Volume :- Non-Ephemeral Storage

Important

Kubectl :- Single cloud.

Kubeadm :- On-premise

Kubefed :- Federated (Hybrid Cloud)

Lec-03 :- K8s Setup Master and Node on Aws.

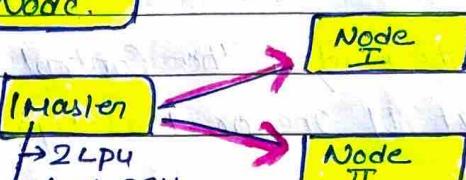
Login into Aws account → Launch 3 instance
⇒ ubuntu 16.04 (t2.Medium)

- ⇒ Now, using puttygen Create private key and save.
- ⇒ Access all the instance (1Master, 2 Nodes)

COMMON COMMANDS FOR Master Node.

```
$ #> sudo su  
$ #> apt-get update.  
$ #> apt-get install apt-transport-https.
```

```
$ #> apt install docker.io -y  
$ #> docker --version.  
$ #> systemctl start docker  
$ #> systemctl enable docker.
```



(This https is needed for intra cluster communication particularly from control plane to individual pods)

Setup Open GPG key this required for intra cluster communication it will be added to source key on this node ie. when k8s sends signed info to our host, it is going to accept those information because this open gpg key is present in the source key.

```
$ #> sudo curl -s https://packages.cloud.google.com/apt.. | sudo apt-key add.
```

```
$ #> vim /etc/apt/sources.list.d/kubernetes.list.
```

↳ deb http://apt.kubernetes.io/ kubernetes-xenial main

⇒ Save and exit.

BOOTSTRAPPING THE MASTER NODE (IN MASTER)

\$#) `kubeadm init`

You will get one long command started from "kubeadm join
172.165.158.6443 .."

Copy this command and save to notepad.

Create both .kube and its parent directories (-p)

\$#) `mkdir -p $HOME/.kube`.

Copy Configuration to kube directory (in config file)

\$#) `Sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`.

Provide permission for config file.

\$#) `Sudo chown $(id -u):$(id -g) $HOME/.kube/config`

Deploy flannel node network for its repository path Flannel is going to place a binary in each node.

\$#) `Sudo kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/documentation/kube-flannel.yaml`.

\$#) `Sudo kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/documentation/k8s-manifests/kube-flannel-rbac.yaml`

*> Configure Worker Node :-> paste long command in both the nodes

*> Go to Master Node :=> \$#) `kubectl get nodes`

(it will provide the Node Status)

Lec-04 - Kubernetes in detail.

Kubernetes Objects.

- ⇒ Kubernetes uses objects to represent the state of the your cluster.
- ⇒ What Containerized applications are running (and on which node)
- ⇒ The policies around how those applications behave, such as restart policies, upgrades and fault tolerance.
- ⇒ Once you create the object, the Kubernetes system will constantly work to ensure that object exist and maintain the cluster's desired state.
- ⇒ Every Kubernetes object includes two nested fields that govern the object Config the Object Spec and the Object Status.
- ⇒ The Spec, which we provide, describes your desired state for the object - the characteristics that you want the object to have.
- ⇒ The Status describes the actual state of the object and is updated and updated by the Kubernetes system.
- ⇒ All objects are identified by a unique name and a UID.

The Basic Kubernetes Objects include.

- 1.7 Pod 2.7 Service 3.7 Volume 4.7 Namespace 6.7 Replicates
- 6.7 Secrets 7.7 ConfigMaps 8.7 Deployment 9.7 Jobs 10.7 Daemonsets.

Relationship between these objects.

- ⇒ pod Manager Containers.
- ⇒ Replicaset Managed pods.
- ⇒ Service expose pod processes to the Outside world.
- ⇒ Configmaps and Secrets helps you configure pods.

Kubernetes Object.

- ⇒ It represents as JSON or YAML files.
- ⇒ You Create them and then push them to the Kubernetes API with kubectl.

State of the Object

- | | |
|-----------------|--------------------------|
| ⇒ Replica (2/2) | ⇒ image (Tomcat /Ubuntu) |
| ⇒ None | ⇒ part |
| ⇒ Volume | ⇒ startup |
| | ⇒ detached (default) |

Kubernetes Objects Management.

The kubectl command line tool supports several different ways to Create and Manage Kubernetes object.

Management Technique	Operate ON	Recommend Env
Imperative Commands	Live Objects	Deployment Project
Declarative object Config	Individual files (yaml/json)	Production

Declarative is about describing what you are trying to achieve, without instructing how to do it.

Imperative, explicitly tells "how to accomplish it"

Fundamental of pods

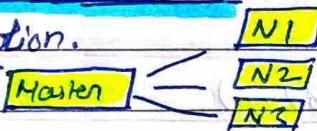
- ⇒ When a pod gets created, it is scheduled to run on a node in your cluster.
- ⇒ The pod remains on that node ~~with~~ until the process is terminated. If the pod object is deleted, the pod is evicted for lack of resource or the node fails.
- ⇒ If a pod is scheduled to a node that fails, or if the scheduling operation itself fails, the pod is deleted.
- ⇒ If a node dies, the pods scheduled to that node are scheduled for deletion after a timeout period.
- ⇒ A given Pod (uid) is not "rescheduled" to a new node, instead it will be replaced by an identical pod. with even the same name if desired, but with a new UID.
- ⇒ Volume in a pod will exist as long as that pod (with that uid) exists. If that pod is deleted for any reason, volume is also destroyed and created again on new pod.
- ⇒ A Controller can Create and Manage multiple pods, handling replication, rollout and providing self-healing capabilities.

Kubernetes Configuration.

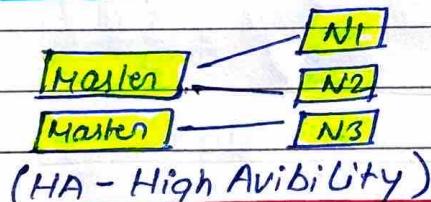
All in One Single node installation.

↳ 2 vcpu (For practice)

Single - Node etcd, Single Master, and Multi - worker installation.



Single node etcd, ~~Multiple~~ Multi Master and Multiworker install



With all in - one, all the master and worker components are installed on a single node. This is very useful for learning, development and testing. This type should not be used in production. Minikube is one such example, and we are going to explore it.

In this setup, we have a single Master node, which also runs a single - node etcd instance. Multiple worker nodes are connected to the Master node.

In this step, we have multiple Master nodes which works in an HA mode, but we have a single - node etcd instance. Multiple worker nodes are connected to the Master node.

Go to Aws account \Rightarrow Launch instance

\Rightarrow Ubuntu 18.04 \Rightarrow t2 medium (2 vcpu)

Now access Fc2 via putty \Rightarrow Login as "ubuntu"

\Rightarrow sudo su

\Rightarrow apt update & apt -y install docker.io

\Rightarrow install kubectl

⇒ install Minikube.

⇒ apt install Contrack

⇒ minikube status.

⇒ kubectl version

⇒ kubectl get nodes

Then check the nodes details :-

⇒ kubectl get nodes. ↳ (output will be like)

o/p	None	Status	Roles	Ages	Version
	ip.172.31.34.55	Ready	Master	2min	v1.9.0-7

⇒ kubectl describe node 172.31.34.35

↳ (Node ip)

⇒ Sample file to write Manifest file for kubernetes.

Kind: Pod

apiVersion: v1

Metadata:

name: testpod

Spec:

Containers:

- name: COO

image: ubuntu

Command: ["/bin/bash", "-c", "while true; do echo Hello-Saurav; sleep 5; done"]

restartPolicy: Never # Defaults to Always.

⇒ kubectl apply -f pod1.yaml (to apply the yaml file)

Multi Container pod Environment

Kind: Pod

apiVersion : v1

Metadata :

name : testpod3

Spec :

Containers :

- name: CO0

image: ubuntu

command : ["/bin/bash", "-c", "while true; do echo Sawan; sleep 5; done"]

- name : CO1

image: ubuntu

command : ["/bin/bash", "-c", "while true; do echo Sawan; sleep 5; done"]

Pod Environment Variables.

Kind: Pod

apiVersion : v1

Metadata :

name : environments

Spec :

Containers :

- name: CO0

image: ubuntu

Command: ["/bin/bash", "-c", "while true; do echo Sawan; sleep 5; done"]

Env: # List of environments .

- name : MYNAME

value: SAWAN

Pod with PORTS

Kind : Pod

apiVersion : V1

Metadata :

name : testpod4

Spec :

Containers :

- name : COO

image :

ports :

- ContainerPort : 80

⇒ Kubectl apply -f pod2.yaml

⇒ Kubectl get pods → (To check the pod and Container details)

⇒ Kubectl describe pod testpod 3 → (describe command to check detail.)

⇒ Kubectl logs -f testpod 3

↳ you will get error, need to specify Container if we need to check the logs of multiple containers.

⇒ Kubectl logs -f testpod 3 -c COO → (Container name)

⇒ Kubectl logs -f testpod 3 -c COO1 → (Container name)

⇒ Kubectl exec testpod3 it -c CO1 - /bin/bash → (Container name)

To login .

Lec - 05 - Kubernetes in detail II

Labels and Selections

- ⇒ Labels are the mechanism you use to organise Kubernetes objects.
(None: Sawan)
- ⇒ A label is a key-value pair without any predefined meaning that can be attached to the objects.
- ⇒ Multiple labels can be added to a single object.
- ⇒ Labels are similar to tags in AWS or git where you use a name to quickly reference.
- ⇒ So you are free to choose labels as you need it to refer on Env which is used for dev or testing or prod, refer a prod group like Department A - Department B

EXAMPLE OF LABELS.

Kind: Pod

apiVersion: v1

Metadata:

name: **Sawanpod**

labels:

env: **development**

class: pods

Spec:

Containers:

- name: **c00**

image: **Ubuntu**

command: **["/bin/bash", "-c", "while true; do echo Hello-Sawan; sleep 5; done"]**

⇒ Multiple labels we can add into a single object.

⇒ Kubectl apply -f pods.yaml

⇒ Kubectl get pods --show-labels
└─(nodes)

↳ if we want to check the labels
we can add node option.

⇒ Now, if you want to add a label to an existing pod.

⇒ Kubectl label pods devlpod Myname = bhuvinesh
└─(object name) └─(object type) └─(key) (Selector)

⇒ Kubectl get pods --show-labels

↳ now, list pods matching a label.

⇒ Kubectl get pods -l env=development.

↳ now, give a list, where 'development' label is not present

⇒ Kubectl get pods -l env!=development

↳ now if you want to delete pod using labels.

⇒ Kubectl delete pod -l env!=development

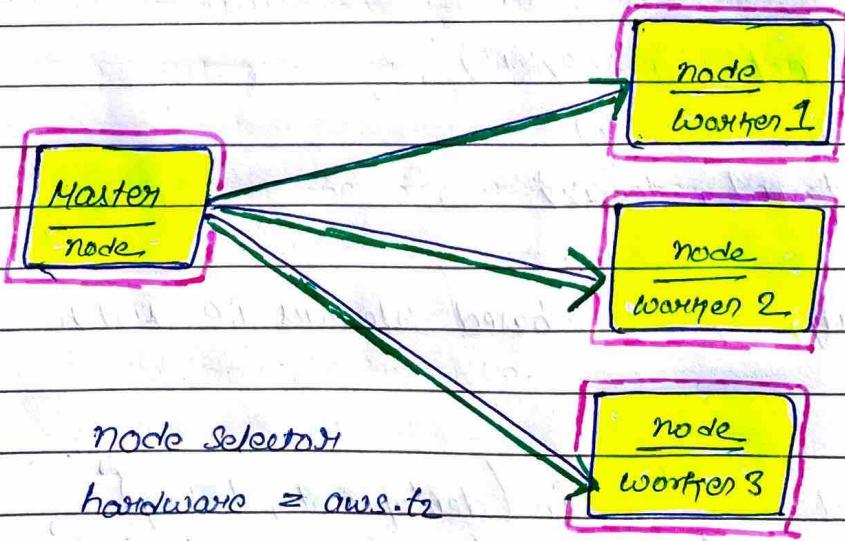
⇒ Kubectl get pods.

Label - Selectors.

- ⇒ Unlike name / UID's labels do not provide uniqueness, as in general, we can expect many objects to carry the same label.
- ⇒ Once labels are attached to an object, we would need filters to narrow down and these are called as label selectors.
 - ⇒ The API currently supports two types of selectors — Equality based and set based.
 - ⇒ A label selector can be made of multiple requirement which are comma-separated.
 - ⇒ Equality Based ($=$, $!=$)
name: sawan
class: Node1
project: development.
 - ⇒ Set Based : (in, notIn and exists)
env: in (production, dev)
env: notIn (team1, team2)
 - ⇒ Kubernetes also supports set-based selectors i.e. match multiple values.
 - ⇒ `kubectl get pods -l 'env:in(development, testing)'`
⇒ `kubectl get pods -l 'env:notIn(development, testing)'`
⇒ `kubectl get pods -l class=prod, myname=bhupinder`

Node - Selector

- 2) One way for selecting labels is to constrain the set of nodes onto which a pod can schedule i.e. you can tell a pod to only be able to run on particular nodes.
- 2) Generally such containers are unnecessary as the scheduler will automatically do a reasonable placement, but in certain circumstances we might need it.
- 2) We can use labels to tag nodes and bind them.
- 2) If the nodes are tagged, you can use the label selector to specify the pods run only of specific nodes.
- 2) First we give label to the node.
- 2) Then we need to use node selector to the pod configuration.



Node Selector Example

Kind: Pod

apiVersion: v1

Metadata:

name: nodeTables

tables:

env: development

Spec:

Containers:

- name: COO

image: ubuntu

command: ["/bin/bash", "-c", "while true; do echo Hello-Sawan;
sleep 5; done"]

nodeSelector:

hardware: t2-medium

Scaling and Replication.

⇒ Kubernetes was designed to orchestrate multiple containers and replication.

⇒ Need for multiple containers / replication helps us with these

Reliability :- By having multiple versions of an application, you prevent problems if one or more fail.

Load Balancing :- Having multiple version of a container enables you to easily send traffic to different instances to prevent overloading of a single instance or node.

Scaling :- When load does become too much for the number of existing instances.

Kubernetes enables you to easily scale up your application adding additional instances as needed.

Rolling update :- Updates to a service by replacing pods one by one.

Replication Controller

- ⇒ A replication controller is an object that enables you to easily create multiple pods, then make sure that number of pods always exists.
- ⇒ If a pod created using RC will be automatically replaced if they die, crash, failed, or terminated.
- ⇒ RC is recommended if you just want to make sure 1 pod is always running, even after system reboots.
- ⇒ You can run the RC with 1 replica & the RC will make sure the pod is always running.

Kind : Replication controller.

Replication Controller

kind: ReplicationController → this defines to create the object of
apiVersion: v1
Replication type

Metadata:

name: myreplica

Spec:

replicas: 2 → this element defines the desired number of pods

Selector: → tells the controller which pods to watch/belong to RC

matchLabels: sawan → this must match the labels.

template: → template element defines a template to launch a new pod

Metadata:

name: testpod0

labels: → Selector value needs to match the label values specified

matchLabels: sawan. in the pod template.

Spec:

Containers:

- name: con

image: ubuntu

command: ["bin/bash", "-c", "while true; do echo hello-sawan; sleep 5; done"]

\$#) kubectl apply -f myrc.yaml

\$#) kubectl get rc

↳ (replication controller)

Name	Desired	Current	Ready	Age
myreplica	5	5	5	48s

\$#) kubectl describe rc myreplica

\$#) kubectl get pods

↳ (name)

\$#> kubectl get pods --show-labels. \Rightarrow To check the labels on Pod

Name	Ready	Status	Replicas	Age	Labels
myreplica-h12	1/1	Running	0	2M40S	myname = Savan
myreplica-t79	1/1	Running	0	2M50S	myname = Savan
myreplica-189	1/1	Running	0	4M40S	myname = Savan
myreplica-xyz	1/1	Running	2	3H10S	myname = Savan

\$#> kubectl scale --replicas=8 rc -l myname=bhupinder.

This will scale the myname = bhupinder template to 8 pods.

\$#> kubectl get pods \Rightarrow to get the pod details.

\$#> kubectl scale --replicas=1 rc -l myname = bhupinder.

\hookrightarrow This command will scale down the pods from 8 to 1.

\$#> kubectl get rc.

\hookrightarrow To check the replication controller details.

Replica Set

\Rightarrow Replica Set is a next generation Replication Controller.

\Rightarrow The Replication Controller only supports Equality based Selection whereas the replica set supports Set-based Selection i.e filtering according to set of values.

\Rightarrow ReplicaSet rather than the Replication Controller is used by other objects like deployment.

Example of Replica Set

Kind : Replicaset

apiVersion : apps/v1

Metadata :

name : myrs

Spec :

replicas : 2

Selector :

MatchExpressions:

- { key : myname, operator : In, values : [Pawan, SAWAN, SAWAN] }

- { key : env, operator : NotIn, values : [production] } # there must match the label.

template :

Metadata :

name : testpod7

Labels :

myname : Pawan

Spec :

Containers :

- name : COO

image : ubuntu

command : ["/bin/bash", "-c", "while true; do echo Hello-Sawan; sleep 5; done"]

\$ #> kubectl apply -f myrs.yaml

\$ #> kubectl get rs .

\$ #> kubectl get pods

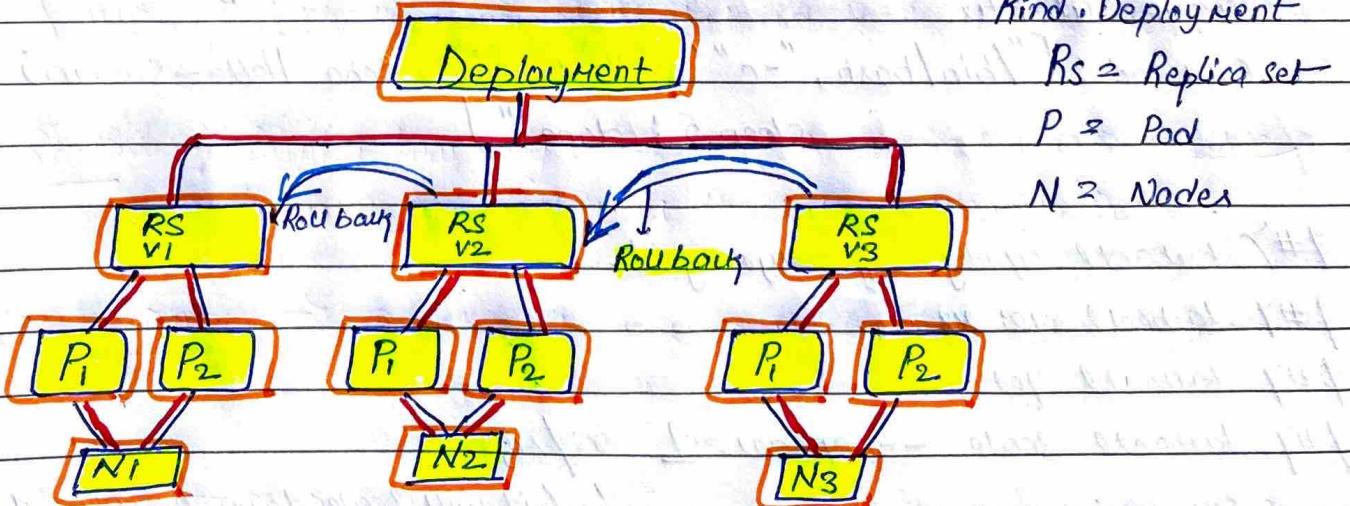
\$ #> kubectl scale --replicas=1 rs/myrs

↳ (it will scale down from 2-1)

Lec :- 06 - Deployment & Rollback

Replication Controller & Replica Set is not able to do update & rollback apps in the cluster.

- ⇒ A deployment object act as a supervisor for pods ; giving you fine - grained control over how and when a new pod is rolled out, updated or rolled back to a previous state.
- ⇒ When using deployment object , we first define the state of the app. then K8s Cluster Schedule Mentioned app instances onto Specific individual Nodes.
- ⇒ K8s then Monitors , if the node hosting an instance goes down or pod is deleted the deployment controller replaces it.
- ⇒ This provides a Self-healing mechanism to address machine failure or maintenance .



- ⇒ A deployment provides declarative update for pod & Replicaset

The following are typical use cases of Deployments :-

- 1.) Create a deployment to rollout a Replicaset :- The replicaset creates pods in the background checks the status of the rollout to see if it succeeded or not.
- 2.) Declare the new state of the pods :- By updating the pod template spec of the deployment. A new replicaset is created and the Deployment Manager moves the pods from the old replicaset to the new one at a controlled rate. Each new replicaset updates the revision of the deployment.
- 3.) Rollback to an earlier deployment Revision :- If the current state of the deployment is not stable each rollback updates the revision of the deployment.
- 4.) Roll back to an earlier deployment Revision :— If the current state of the deployment is not stable each rollback updates the revision of the deployment.
- 5.) Scale up the deployment to facilitate more load.
- 6.) Pause the deployment to apply multiple fixes to its pod template spec and then resume it to start a new rollout.
- 7.) Cleanup older replicsets that you don't need anymore.

⇒ if there are problems in the deployment, k8s will automatically rollback to the previous version, however you can also explicitly rollback to a specific revision, or in our case to Revision 1 (the original pod version)

⇒ you can rollback to a specific version by specifying it with --to-revision.

#⇒ For eg ⇒ kubectl rollout undo Previous Version.

Type → deploy/mydeployment --to-revision=2

Note ⇒ That the name of the Replicaset is always formatted as
[Deployment-name]
[Random]

CMD ⇒ kubectl get deploy.

* When you inspect the deployment in your cluster, the following field are displayed:

NAME ⇒ List the names of the deployments in the namespace.

READY ⇒ Display how many replicas of the application are available to your user if follows the pattern ready / desired.

UP-TO-DATE ⇒ Display the number of replicas that have been updated to achieve the desired state.

AVAILABLE ⇒ Displays how many replicas of the application are available to your user.

AGE ⇒ Display the amount of time that the Application has been Running.

Lab Login to AWS Management Console, Create one Ubuntu(t2-medium) instance. Take access using putty.

\$#> Sudo su

\$#> Sudo apt update && apt -y install docker.io

\$#> install kubectl

\$#> install Minicube also

\$#> apt install Contrack

\$#> minicube start --VM-drivers=none

\$#> minicube status

\$#> Kubectl version

\$#> Kubectl get nodes

Deployment YAML Example

Kind : Deployment

apiVersion : apps/v1

Metadata :

name : mydeployment

Spec :

Replicas : 2

Selector :

MatchLabels :

name : deployment

Template :

Metadata :

name : testpod

Labels :

name : deployment

Spec :

Containers :

- name : coo

image : ubuntu

command : ["/bin/bash", "-c", "while true; do echo Hello-sawan; sleep 5; done"]

⇒ To check deployment was created or not.

\$#> kubectl get deploy.

⇒ To check how deployment Created RS & pods

\$#> kubectl describe deploy mydeployments

\$#> kubectl get rs

⇒ To scale up or Scale down

\$#> kubectl scale --replica=1 deploy mydeployments.

Object type ↳ Object name

To check, what is running inside container.

\$#> kubectl logs -f < podname >

\$#> kubectl rollout status deployment mydeployments

\$#> kubectl rollout history deployment mydeployments

\$#> kubectl rollout undo deployment mydeployments.

Failed Deployment your deployment may get stuck trying to deploy

its newest Replicaset without ever completing.

This can occur due to some following reasons.

i) Insufficient Quotas

ii) Readiness probe failures.

iii) Image pull errors

iv) insufficient permission

v) Limit Range

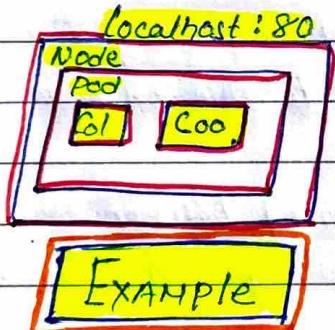
vi) Application runtime Misconfig.

Lec-07 - Kubernetes Networking, Services, Nodeport, Volumes

Kubernetes Networking addresses four Concerns :

- i) Containers within a pod use networking to communicate via Loopback.
- ii) Cluster Networking provides communication between different pods.
- iii) The Service resource lets you expose an application running in pods to be reachable from outside your cluster.
- iv) You can also use services to publish services only for consumption inside your cluster.

Cluster : Container to Container communication on same pod happens through localhost within the container.



Kind : Pod
apiVersion : v1

Metadata :
name : testpod

Spec :

Containers :

- name : coo

image : ubuntu

command : ["/bin/bash", "-c", "while true; do echo Hello - sawan; sleep 5; done"]

- name : co1

image : httpd

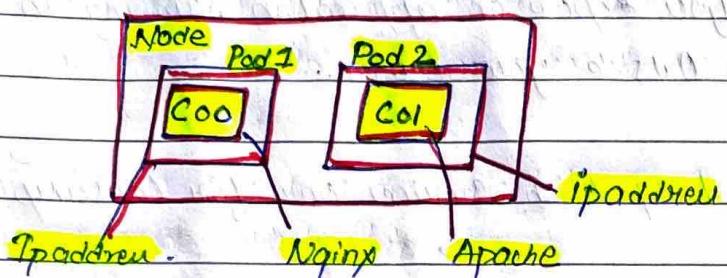
Ports :

- ContainerPort : 80

Now try to establish communication between two different pods within some node.

⇒ Pod to pod communication on worker nodes happen through pods

⇒ By default pod IP will not be accessible outside the node.



EXAMPLE FOR Nginx

Kind : Pod

apiVersion : v1

Metadata :

name : testpod1

Spec :

Containers :

- name : col1

image : nginx

ports :

- ContainerPort : 80

\$ #> kubectl apply -f pod2.yml

Example for Apache

Kind : Pod

apiVersion : v1

Metadata :

name : testpod4

Spec :

Containers :

- name : l03

image : httpd

Ports :

- ContainerPort : 80

\$#> kubectl apply -f pod3.yaml.

\$#> kubectl get pods -o wide #> For all details like ip address

\$#> curl 172.*.*.1 → pod 1 ip address

\$#> curl 172.*.*.4 → pod 2 ip address.

Object - Service :

Each pod gets its own ip address, however in a deployment, the set of pods running in one moment in the time could be different from the set of pods running that application a moment later.

This leads to a problem : if some sets of pods (call them 'backend') provide functionality to other pods (call them frontend) inside your cluster, how do the frontends find out and keep track of which ip address. To connect, so that the frontend can use the backend part of the workload.

- ⇒ When using RC, pods are terminated and created during scaling or replication operations.
- ⇒ When using deployments, while updating the image version the pods are terminated and new pods take the place of other pods.
- ⇒ pods are very dynamic i.e. they come & go on the k8s cluster and on any of the available nodes & it would be difficult to access the pods as the pod ip changes once it's recreated.
- ⇒ Service object is an logical bridge between pods and end users, which provides virtual ip.
- ⇒ Services allows clients to reliably connect to the containers running in the pod using the VIP
- ⇒ The VIP is not an actual IP connected to a network interface, but its purpose is purely to forward traffic to one or more pods.
- ⇒ Kube proxy is the one which keeps the mapping between the VIP and the pods up-to-date. which queries the API Server to learn about new services in the cluster.
- ⇒ Although each pod has a unique IP address, those IPs are not exposed outside the cluster.
- ⇒ Services helps to expose the VIP mapped to the pods & allows application to receive traffic.

- ⇒ Labels are used to select which are the pods to be put under a service.
- ⇒ Creating a Service will create an endpoint to access the pods / application in it.
- ⇒ Services can be exposed in different ways by specifying a type in the Service Spec:

#⇒ Cluster IP

#⇒ NodePort

#⇒ Load Balancer

Loadbalancer :- Created by cloud providers that will route external traffic to every node on the Nodeport. (Eq F1B on AWS)

Headless :- Creates several endpoints that are used to produce DNS Records. Each DNS record is bound to a pod.

- ⇒ By default Service can run only between ports 30,000 - 32767
- ⇒ The set of pods targeted by a service is usually determined by a selector.

Example of Deployment

Kind: Deployment

apiVersion: apps/v1

Metadata: name of deployment do stored this name in pod

Name: mydeployments

Spec:

Replicas: 1

Selector:

MatchLabels:

name: deployment

template:

Metadata:

name: testpod1

Labels:

name: deployment

Spec:

Containers:

- name: c00

image: httpd

ports:

- containerPort: 80

\$#> kubectl apply -f deploymenthttpd.yaml

\$#> kubectl get pods

\$#> kubectl get pods -o wide

\$#> curl 172.16.12.* → pod IP to check httpd is working or not.

Example of Service

Kind : Service # Defines to Create Service type Object

Service

apiVersion : v1

Metadata :

name : demo service

Spec :

Ports :

- port : 80 # Container port exposed.

targetPort : 80 # Pod Port

Selector :

name : deployment # Apply this service to any pod which has the

type : ClusterIP

specify label

→ specifies the service type i.e. Cluster IP or Nodeport

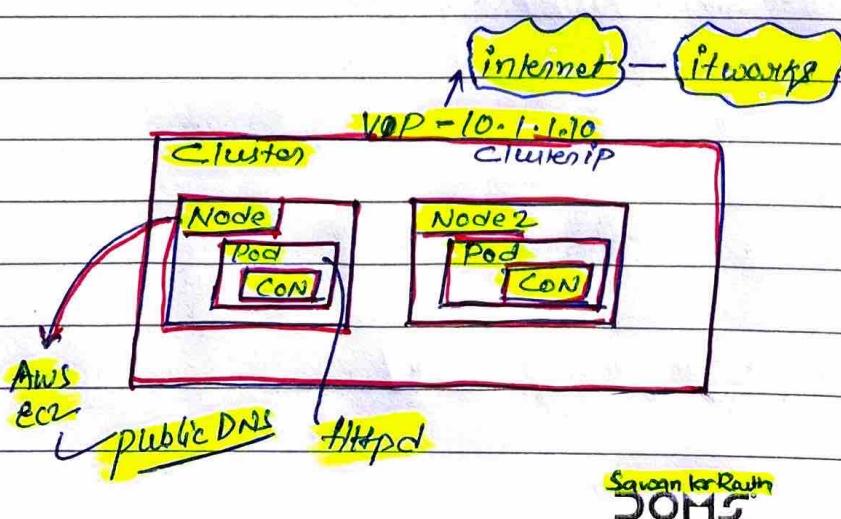
\$#> kubectl apply -f service.yaml

\$#> kubectl get svc

\$#> curl 10.231.10.* ⇒ cluster ip

NODEPORT :- ⇒ makes a service accessible from outside the cluster.

⇒ Expose the service on the same port of each selected node
int the using NAT



Example of NODEPORT

Kind: Service # Defines to create service type object

apiVersion: v1

Metadata:

name: demoservice

Spec:

ports:

- port: 80 # Container port exposed.

targetPort: 80 # Pod port (internal)

Selector:

name: deployment # Apply this service to any pod which has the

type: Nodeport

Specific label.

Specifies the service type is ClusterIP or

Nodeport

\$#) kubectl apply -f deployment.yaml

\$#) kubectl apply -f svc.yaml

\$#) kubectl get svc

it will show the Nodeport for example

80:31341/TCP

\$#) kubectl describe svc demoservice

K8s Volumes

- ⇒ Containers are short lived in Nature.
- ⇒ All data stored inside a container is deleted if the container Crashes. However the kubelet will restart it with a clean state, which means that will not have any of the old data.
- ⇒ To overcome this problem, K8s uses volume. A volume is essentially directory backed by a storage medium. The storage medium and its content are determined by the volume type.
- ⇒ In K8s, a volume is attached to a pod and shared among the containers of that pod.
- ⇒ The volume has the same life span as the Pod, and its outlives the containers of the pod. This allows data to be preserved across container restarts.

K8s Volume Types

- ⇒ A volume type decides the properties of the directory, like size, content etc. Some examples of volume types are :-
- ⇒ Node-local types such as `emptydir` and `hostpath`.
- ⇒ File sharing types such as `nfs`.
- ⇒ Cloud provider-specific types like AWS `EBS`, Azure disk.
- ⇒ Distributed file system types, for example `glusterfs` or `cephfs`.
- ⇒ Special purpose types like `secret`, `git repo`.

Empty Dir. :-

↳ Volume Type

- ⇒ Use this when we want to share contents between multiple containers on the same pod & not to the host machine.
- ⇒ An EmptyDir volume is first created when a pod is assigned to a node and exits as long as the pod is running on that node.
- ⇒ As the name says, it is initially empty.
- ⇒ Containers in the pod can all read and write the same files in the EmptyDir volume, through that volume can be mounted at the same or different paths in each container.
- ⇒ When a pod is removed from a node for any reason, the data in the empty dir is deleted forever.
- ⇒ Container crashing does not remove a pod from a node, so the data in an empty dir volume is safe across container crashes.

Example of EmptyDir

apiVersion: v1

kind: Pod

Metadata:

name: myvolemptydir

Spec:

Containers:

- name: c1

Image: Centos

Command: ["bin/bash", "-c", "sleep 5000"]

Volume Mounts:

- name : xchange

MountPath : "/tmp/xchange"

- name : c2

image : Centos

command : ["/bin/bash", "-c", "sleep 1000"]

volumeMounts :

- name : xchange

Mountpath : "/tmp/data"

Volume :

- name : xchange

emptyDir: {}

\$#) kubectl apply -f emptydir.yaml

\$#) kubectl get pods

\$#) kubectl exec myvolumeemptydir -c c1 -it -- /bin/bash
↳ pod name ↳ Container

\$#) cd /tmp

\$#) ls ⇒ you will find the volume name.

\$#) then go to the volume value and create some txt file.

\$#) Then go to C2 ⇒ Container 2

\$#) kubectl exec myvolumeemptydir -c c2 -it -- /bin/bash

\$#) cd /tmp

\$#) ls ⇒ All files will be also here.

Host Path :-

⇒ we use this when we want to access the Content of a pod / container from hostnames.

⇒ A hostpath Volume Mounts a file or directory from the host node's filesystem into your pod.

Example of Hostpath

apiVersion: v1

Kind: Pod

Metadata:

name: volumehostpath

Spec:

Containers:

- image: Centos

name: testc

command: ["/bin/bash", "-c", "sleep 15000"]

Volume Mounts:

- Mount Path: /tmp/hostpath

name: testvolume

Volumes:

- name: testvolume

hostPath:

path: /tmp/data

\$#) kubectl apply -f hostpath.yaml

\$#) kubectl get pods.

\$#) kubectl exec volumehostpath -- ls /tmp

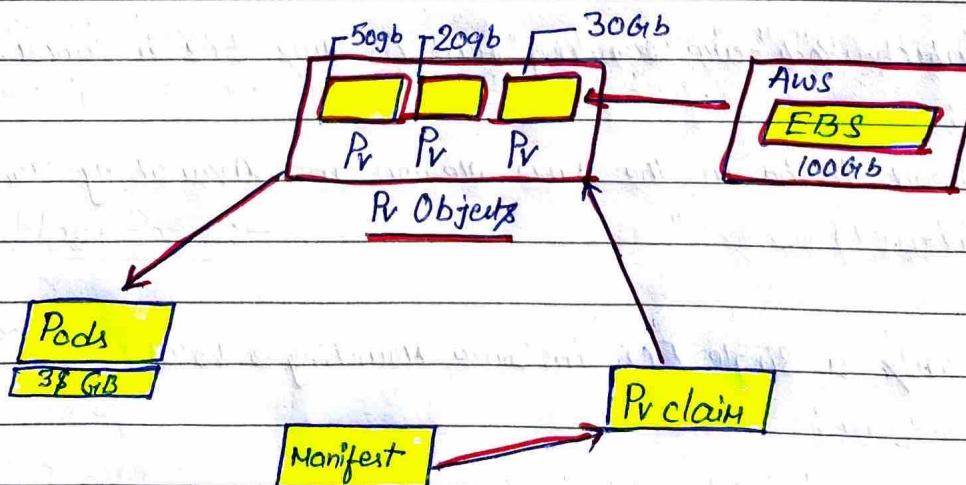
\$#) echo -e "I am sawan" >myfile

\$#) kubectl exec volumehostpath -- ls /tmp/hostpath

\$#) kubectl exec myvolhostpath -- cat /tmp/hostpath/myfile

Lec-08 - Persistent Volume & Liveness Probe

- ⇒ In a typical IT env, storage is managed by the storage/system admin. The end user will just get instructions to use the storage, but does not have to worry about the underlying storage management.
- ⇒ In the containerized world, we would like to follow similar rules, but it becomes challenging, given the many volume types we have seen earlier. Kubernetes handles this problem with the persistent Volume subsystem.
- ⇒ A persistent Volume is a cluster-wide resource that you can use to store data in a way that is persist beyond the life-time of a pod.
- ⇒ The pv is not backed by locally attached storage on a worker node but by networked storage system such as EBS or NFS or a distributed filesystem like ceph.
- ⇒ K8s provides API for users and admin to manage and consume storage to manage the volume, it uses the persistent volume API resource type and to consume it, uses the persistent volume-claim API resource type.



Persistent Volume Claim

- ⇒ In order to use a PV you need to claim it first, using a persistent volume claim.
- ⇒ The PVC request a PV with your desired specification (size, access modes, speed etc) from k8s and once a suitable Persistent Volume is found, it is bound to a persistent volume claim.
- ⇒ After a successful bind to a pod, you can mount it as a volume.
- ⇒ Once it is user finished its wants, the attached PV can be released then underlying PV can then be reclaimed and recycled for future usage.

Aws EBS

- ⇒ An aws ebs volume mounts an Aws EBS volume into your pod unlike EmptyDir, which is erased when a pod is removed; the contents of an EBS volume are preserved and the volume is merely unmounted.

There are some restrictions:-

- i) The nodes on which pods are running must be Aws Ec2 instances.
- ii) Those instances need to be in the same region and availability zone as the EBS volume.
- iii) EBS only supports a single Ec2 instance mounting a volume.

Example of Persistent Volume

apiVersion : v1

kind : PersistentVolume

Metadata :

name : mypvcvol

Spec :

capacity :

storage : 1Gi

accessMode :

- ReadWriteOnce

PersistentVolumeReclaimPolicy : Recycle

awsElasticBlockStorage :

volumeID : vol-0dc02983B1XYZ # Put EBS volume ID Here

FsType : ext4

\$#> kubectl apply -f mypv.yaml

\$#> kubectl get pv

Example of Persistent Volume Claim

apiVersion : v1

kind : PersistentVolumeClaim

Metadata :

name : mypvcvolumeclaim

spec :

accessModes :

- ReadWriteOnce

resources :

requests :

storage : 1Gi

\$#) kubectl apply -f mypv.yaml
\$#) kubectl get pvc

Example of Deployment.

apiVersion: apps/v1

kind: Deployment

metadata:

name: mydeploy

spec:

replicas: 1

selector: # tells the controller which pods to watch / belong to

matchLabels:

app: mypv

template:

metadata:

labels:

app: mypv

spec:

containers:

- name: shell

image: Centos

command: ["/bin/bash", "-c", "sleep 1000"]

volumes:

- name: mypd

mountPath: "/tmp/persistent"

volumes:

- name: mypd

persistentVolumeClaim:

claimName: mypvc.volclaim

```
$#> kubectl apply -f deployrc.yaml  
$#> kubectl get deploy  
$#> kubectl get rs  
$#> kubectl get pods  
$#> kubectl exec pvdeploy-0589xyz -it -- /bin/bash  
$#> cd /tmp/persistent
```

Now delete the existing pod and create an RS will create one new pod to Under that pod in the tmp dir and check the same file will be there.

Health Check / Liveness Probe

- ⇒ A pod is considered ready when all of its containers are ready.
- ⇒ In order to verify if a container in a pod is healthy and ready to serve traffic, k8s provides for a range of healthy checking connection.
- ⇒ Health checks or probes are carried out by the kubelet to determine when to ~~restart~~ ^{recreate} a container (For liveness probe) and used by services and deployments to determine if a pod should receive traffic.

For eg. Liveness probes could catch a deadlock, where an application is running but unable to make ~~any~~ progress. Restoring - Restarting a container in such a state can help to make the application make available despite bugs.

- ⇒ One use of readiness probe is to control which pods are used as backends for services when a pod is not ready, it is removed from Service Load Balancer.
- ⇒ For running healthcheck, we would use cmd specific to the application
- ⇒ if the cmd succeeds, it returns 0, and kubelet considers the container to be alive and healthy. if the command returns a non-zero value, the kubelet kills the pod and removes it.

Example of liveness Probe

Kind : Pod

Metadata :

Labels :

test : liveness

name : mylivenessprobe

Spec :

Containers :

- name : liveness

image : ubuntu

args :

- /bin/sh

--c

- touch /tmp/healthy ; sleep 1000

livenessprobe :

define health check

exec :

command :

command to run periodically

- cat :

- /tmp/healthy

initialDelaySeconds: 5 # Wait for the specified time before it runs first probe.
periodSeconds: 5
timeoutSeconds: 30 # Run the above command every 5 sec

\$#> kubectl apply -f liveprobe.yaml

\$#> kubectl get pods.

\$#> kubectl describe pod myliveprobe.

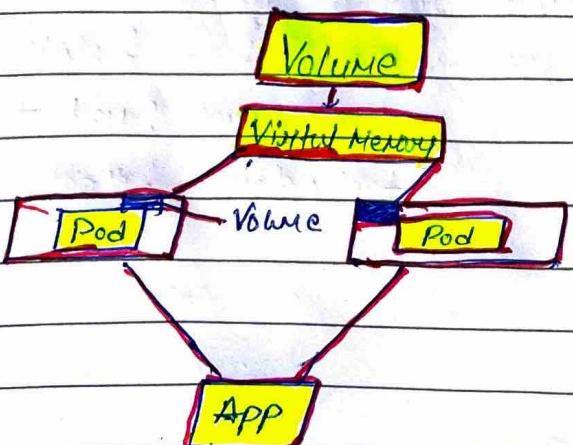
\$#> kubectl exec myliveprobe -it -/bin/bash.

\$#> cat /tmp/healthy.

\$#> echo \$? # to check the status.

Lec -09 - ConfigMap and Secrets

- ⇒ while performing application deployments on k8s cluster, sometimes we need to change the application Configuration file depending on environment like dev, QA, stage or prod.
- ⇒ Changing this application Configuration file means we need to change source code, Commit the change, creating a new image and then go through the completed deployment process.
- ⇒ Hence these Configuration should be decoupled from image Content in Order to keep Containerised application portable.
- ⇒ This is where k8s Configmap Come handy it allows us to handle Configuration files much more efficiently.
- ⇒ Configmaps are useful for storing and sharing non-sensitive; unencrypted Configuration information we Secretly otherwise.
- ⇒ Configmap can be used to store fine-grained information like individual properties or entire config files.
- ⇒ Configmap are not intended to act as a replacement for properties file.



⇒ Configmap can be created in following ways :-

i) As Env Variable.

ii) As volume in the pod.

Kubectl create configmap <mapname> --from-file = <file to read>

SECRETS

⇒ You don't want sensitive information such as a database password or an API key kept around in clear text.

⇒ Secrets provide you with a mechanism to use such information in a safe and reliable way with the following properties :-

⇒ Secrets are namespace objects, that exist in the context of namespace

⇒ you can access them via a volume or an environment variable from a container running in a pod.

⇒ The secret data on node is stored in tmpfs volume (tmpfs is a file system which keeps all files in virtual memory everything in tmpfs is temporary in the sense that no files will be created on your hdd).

⇒ A per-secret size limit of 1mb exist.

⇒ The API Server stores secrets as plaintext in etcd.

Secrets can be created.

1. From a text file

2. From a YAML file.

```
$#> vi sample.conf # create sample config file.  
$#> kubectl create configmap mymap --from-file=sample.conf  
$#> kubectl get configmap # to check the details under the file.  
$#> kubectl describe ConfigMap mymap.
```

Example of ConfigMap

apiVersion: v1

kind: Pod

Metadata:

name: myvolconfig

Spec:

containers:

- name: c1

image: CentOS

command: ["/bin/bash", "-c", "while true; do echo 'Hello Sawan'; sleep 5; done"]

volumeMounts:

- name: testconfigmap

mountPath: "/tmp/config" # the config files will be mounted

Volumes:

- name: testconfigmap

ConfigMap:

name: mymap # This should match the Config Map name created

items:

- key: sample.conf

path: sample.conf

```
$#> kubectl apply -f deployconfigmap.yaml  
$#> kubectl get pods.  
$#> kubectl exec myvolconfig -it -- /bin/bash
```

Example of MYENV

apiVersion: v1

kind: Pod

Metadata:

name: MyEnvConfig

Spec:

Containers:

- name: C1

image: CentOS

command: ["/bin/bash", "-c", "while true; do echo Hello Saman; sleep 5; done"]

Env:

- name: MYENV # env name in which value of the key is stored.

valueFrom:

configMapKeyRef:

name: mymap # name of the config created.

key: sample.conf

```
$#> kubectl apply -f deployenv.yaml
```

```
$#> kubectl get pods.
```

```
$#> kubectl exec myenvconfig -it -- /bin/bash
```

```
$#> env
```

```
$#> echo $MYENV
```

```
$#> exit.
```

Example of Secret

```
$#> echo "Mont" > username.txt ; echo "mypassword" > password.txt  
$#> cat password.txt  
$#> kubectl create secret generic mysecret --from-file=username.txt --from-file=password.txt
```

apiVersion: v1

kind: Pod

Metadata:

name: myvolsecret

Spec:

containers:

-name: c1

image: centos

command: ["/bin/bash", "-c", "while true; do echo Hello-Saavn; sleep 5; done"]

volumeMounts:

-name: testsecret

mountPath: "/tmp/mysecret" # the secret file will be mounted

volumes:

as readonly by default here

-name: testsecret

Secret:

secretname: mysecret.

```
$#> kubectl apply -f deployment.yaml
```

```
$#> kubectl get pods
```

```
$#> kubectl exec myvolsecret -it -- /bin/bash
```

```
$#> cd /tmp
```

```
$#> cat password.txt
```

```
$#> cat username.txt
```

Lec-10 - Namespaces, Limits & Request

Namespaces:- you can name your object, but if many are using the cluster then it would be difficult for managing.

⇒ A namespace is a group of related elements that each have a unique name or identifier. Namespace is used to uniquely identify one or more names from other similar names of different objects & groups in the namespace in general.

⇒ Kubernetes names help different projects teams or customers to share a Kubernetes cluster & provider :-

⇒ A scope for every name.

⇒ A mechanism to attach authorization and policy to a subsection of the cluster.

⇒ By default, a k8s cluster will instantiate its default namespace when provisioning the cluster to hold the default set of pods, services and deployments used by the client cluster.

⇒ We can use resource Quota on specifying how many resources each namespace can use.

⇒ Most k8s resources (e.g. pods, service replication controllers and others) are in some namespaces and low-level resources such as nodes and persistent volumes, are not in any namespace.

⇒ Namespaces are intended for use in environment with many users spread across multiple teams, or projects for clusters with a few to tens of users, you should not need to create and think about namespaces and all.

Create New Namespace

⇒ Let us assume we have shared K8s cluster for Dev and production we care.

⇒ The dev team would like to maintain a space in the cluster where they can get a view on the list of Pod, Service and Deployment they use to build and run their application. For this, no restrictions are put on who can or cannot modify resources to enable agile development.

⇒ For prod team we can enforce strict procedure on who can or cannot manipulate the set of pods, services and deployments.

cmd :- **kubectl get namespaces**.

Namespace YAML

apiVersion : v1

kind : Namespace

Metadata:

name: dev

Labels:

name: dev

\$#> kubectl config set-context \$(kubectl config current-context) --namespace=dev

Creating Sample Pod YAML

Kind : Pod

apiVersion : v1

Metadata :

name : testpod

Spec :

Containers :

- name : pod

image : ubuntu

command : ["/bin/bash", "-c", "while true; do echo Technical gulfqui;
sleep 5; done"]

Restart Policy : Never.

\$#> kubectl apply -f pod.yaml -n dev → it will create the pod under
dev namespace.

\$#> kubectl get pods -n dev

{#> To change the pod view from default namespace to dev namespace
Command is as follow.

\$#> kubectl config set-context \$(kubectl config current-context)
--namespace=dev

{#> To check which is the default namespace.

\$#> kubectl config view | grep namespace:

\$#> kubectl get pods

Managing Compute Resources for Containers

- ⇒ A pod in k8s will run with no limits on CPU and Memory.
- ⇒ You can optionally specify how much CPU and Memory (RAM) each container needs.
- ⇒ Scheduler decides about which node to place pods; Only if the Node has enough CPU resources available to satisfy the pod CPU request.
- ⇒ CPU is specified in units of cores and Memory is specified in units of Bytes.

Two types of ~~Container~~ Constraints can be set for each resource type:

Request and Limit

- ⇒ A request is the amount of that resources that the system will guarantee for the container and kubernetes will use this value to decide on which node to place the pod.
- ⇒ A limit is the max amount of resources that k8s will allow the container to use. In the case that request is not set for a container, it default to limit if limit is not set, then if default to 0.
- ⇒ CPU values are specified in 'MilliCPU' and Memory in MiB.

Resource Quota

⇒ A k8s cluster can be divided into namespaces if a Container is created in a namespace that has a default CPU limit and the container does not specify its own CPU limit, then the Container is assigned the default CPU limit.

⇒ Namespaces can be assigned Resource Quota objects, this will limit the amount of usage allowed to the objects in that namespace.

You can limit:

1. Compute

2. Memory

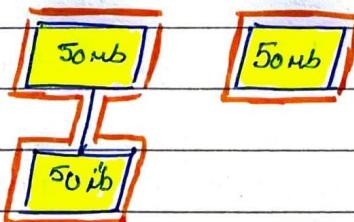
3. Storage

Here are the restrictions that a Resource Quota imposes on a namespace.

⇒ Every Container that runs in the namespace must have its own CPU limit.

⇒ The total amount of CPU used by all containers in the namespace must not exceed a specified limit.

Namespace → 200 mB → RAM



Example of Resource

apiVersion: v1

kind:

Metadata:

name: resource1

Spec:

containers:

- name: resource1

image: Centos

command: ["/bin/bash", "-c", "while true; do echo Hello-Saumon
; sleep 5 ; done"]

resources:

requests:

memory: "64Mi"

CPU: "100M"

limits:

memory: "128Mi"

CPU: "200M"

\$#> kubectl apply -f resource.yaml

\$#> kubectl describe pod resource1

Example of Resource Quota

apiVersion : v1

kind: ResourceQuota

Metadata:

name : myquota

Spec:

hard:

limits.cpu : "400M"

limits.memory : "400Mi"

requests.cpu : "200M"

requests.memory : "200Mi"

\$> kubectl apply -f resourcequota.yaml

Example of Testpod and Replicaset

kind: Deployment

apiVersion: apps/v1

Metadata:

name: deployments

Spec:

replicas: 3

selector:

matchLabels:

objtype: deployment

template:

Metadata:

name: testpods

labels:

objtype: deployment

Spec:

Containers:

- name: COO

Image: ubuntu

command: ["/bin/bash", "-c", "while true; do echo Hello - Savan ; sleep 5; done"]

Resources:

requests:

cpu: "200m"

→ Reduce it to 150m

\$#> kubectl apply -f testpod.yaml

\$#> kubectl get deploy

Lec - 11 - Resource Quota And Horizontal Scaling

pod-nodespace

limits: Cpu : "600m"

limits: Memory : "800Mi"

Cpu limit → 500 m

Request → 72

Request = Limit

requests: Cpu : "200m"

requests: Memory : "500M"

Cpu Request = 600m

Limit = 1

Limit ≠ Request

Default Range

CPU

(min) request = 0.5 CPU

(max) limit = 1

Memory :

(min) request = 500 M

(max) limit = 1G

Example for Memory

apiVersion: v1

kind: LimitRange

Metadata:

name: Mem-min-max-demo-1r

Spec:

Limits:

- Max:

Memory: 1Gi

Min:

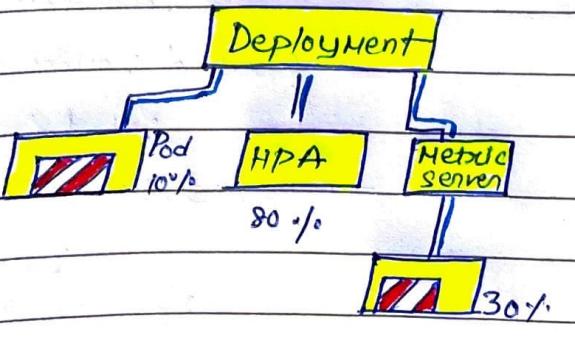
Memory: 500 Mi

Type: Container

\$#) kubectl describe pod podname
\$#) kubectl apply -f yml file name.

Horizontal pod AutoScaler :-

- ⇒ Kubernetes has the possibility to automatically scale pods based on observed CPU utilization, which is Horizontal pod Autoscaling.
- ⇒ Scaling can be done only for scalable objects like controller, deployment or replica set.
- ⇒ HPA is implemented as a k8s API Resource and a Controller.
- ⇒ The controller periodically adjusts the number of replicas in a replication controller or deployment to match the observed average CPU utilization to the target specified by user.
- ⇒ The HPA is implemented as a controlled loop with a period controlled by the controller manager - horizontal - pod - auto - scaler - sync - period flag (Default value of 30 sec).
- ⇒ During each period, the controller manager queries the resource utilization against the metrics specified in each horizontal pod AutoScaler definition —



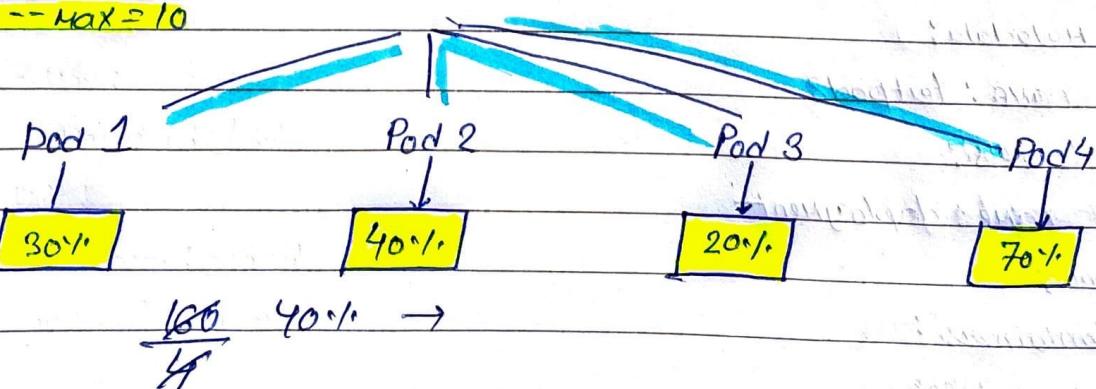
Horizontal Pod Autoscaler (HPA)

- ⇒ For per-pod resource metrics (like CPU), the controller fetches the metric from the Resource Metrics API for each pod targeted by the Horizontal pod Autoscalers.
- ⇒ Then if a target utilization value is set the controller calculates the utilization value as a percentage of the equivalent resource request on the containers in each pod.
- ⇒ If a target raw-value is set, the raw metric values are used directly. The controller then takes the mean of the utilization on the raw value across all targeted pods, and produce a ratio up to scale the number of desired replicas.
- ⇒Cooldown period to wait before another down scale operation can be performed is controlled by `--horizontal-pod-autoscaler-downtime-stabilization` flag (default value of 5 min)
- ⇒ Metric Server needs to be deployed in the cluster to provide metrics via the Resource Metrics API

`--kubernetes-ingress-tr`

\$⇒ Kubecli autoscale deployment mydeploy --cpu-percent=20 --min=1

`--max=10`



Command to install Metric server on k8s.

```
$#) wget -O metricserver.yaml https://github.com/kubernetes-sigs/metric-server/releases/latest/download/components.yaml
```

it will download the metricserver.yaml file, then open that file under vi tool and go under MetricServer line and add the line after

```
-- secure-port = 4443
```

```
-- kubelet-secure-tls # add this line
```

```
$#) kubectl apply -f metricserver.yaml
```

Then Create a Deployment and a Container to watch Realtime scaling.

Kind: Deployment

apiVersion: apps/v1

Metadata:

name: mydeployment

Spec:

replicas: 1

selector:

matchLabels:

name: deployment

template:

Metadata:

name: testpod8

labels:

name: deployment

spec:

containers:

- name: COO

image: httpd

ports:

- containerport: 80

Metadata:

Unit:

CPU: 500m

Requests:

CPU: 200m

\$#> kubectl apply -f deploy.yaml

\$#> kubectl get all

Command to start autoscale on deployment :-

\$#> kubectl autoscale deployment mydeploy --cpu-percent=20
--min=1 --max=10

Now → Enter into the pod and give some load.

\$#> kubectl exec mydeploy-7899fccxyz -it -- /bin/bash
\$#> apt-get update

Now open that terminal ↗ of Kubernetes server.

\$#> watch kubectl get all

it will check and scale automatically in every 2 sec
and create the pod according to the usage.

Lec-12- Jobs, Init Containers & Pod Lifecycle

Jobs :-

→ we have replicsets, Daemonsets, statefleets and deployments they all share one common property : they ensure that their pods are always running if a pod fails, the controller restarts it or reschedules it to another node to make sure the application pods is hosting keeps running.

Use Cases :-

1. Take Backup of a DB
2. Helm charts uses jobs.
3. Running Batch process
4. Run a task at an schedule interval
5. Log rotation.



Example of Batch

apiVersion : batch/v1

kind : Job

Metadata :

name : testjob

Spec :

template :

Metadata :

name : testjob

Spec :

Containers :

- name : counter

image: Centos:7
command: ["/bin/bash", "-c", "echo Hello-sawan; sleep 5"]
restartPolicy: Never

\$#> kubectl apply -f job.yaml
\$#> kubectl get pods
\$#> watch !!

apiVersion: batch/v1
kind: Job
Metadata:
name: testjob
Spec:
parallelism: 5 # Runs four Pod in parallel
activeDeadlineSeconds: 10 # Timeout after 30 sec
template:
Metadata:
name: testjob
Spec:
containers:
- name: Counter
image: Centos:7
command: ["/bin/bash", "-c", "echo Technical - Gruftgu; sleep 20"]
restartPolicy: Never

\$#> kubectl apply -f job2.yaml
\$#> kubectl get pods
\$#> watch !!
\$#> watch kubectl get pods

The "Cron job pattern":

- ⇒ if we have multiple nodes hosting the application for high availability which node handles Cron job?
- ⇒ what happens if multiple identical cron jobs run simultaneously?

Example of Cron

apiVersion: batch/v1beta1

kind: CronJob

Metadata:

name: bhupi

Spec:

Schedule: * * * * *

JobTemplate:

spec:

Containers:

-image: ubuntu

name: sawan

command: ["/bin/bash", "-c", "Hello sawan ; sleep 10;"]

restartPolicy: Never

\$#> kubectl apply -f cron.yaml

\$#> kubectl get pods

\$#> watch !!

\$#> watch kubectl get pods

Init Containers :-

- ⇒ Init Containers are specialised containers that run before app. containers in a pod.
- ⇒ Init containers always runs to completion.
- ⇒ If a pod's init container fails Kubernetes repeatedly restarts the pod until the init container succeeds.
- ⇒ init containers do not support readiness probe.

Use Cases :-

- ⇒ Seeding a Database.
- ⇒ Delaying the application launch until the dependencies are ready.
- ⇒ Clone a git repo into a volume.
- ⇒ Generate configuration file dynamically.

Example of Tnit

apiVersion : v1

kind: Pod

Metadata:

name: initContainer

Spec:

initContainers:

- name: c1

image: centos

command: ["/bin/sh", "-c", "echo Helloall >/tmp/xchange/testfile;

sleep 5;"]

volumeMounts:

- name: xchange

mountPath: "/tmp/xchange"

Containers:

- name: c2

image: centos

command: ["/bin/bash", "-c", "while true; do echo 'cat /tmp/data/

terfile; sleep 5; done"]

volumeMounts:

- name: xchange

mountPath: "/tmp/data"

volumes:

- name: xchange

emptyDir: {}

\$#> kubectl apply -f init.yaml

\$#> kubectl get pods

\$#> watch !!

\$#> watch kubectl get pods

\$#> kubectl logs -f pod/initcontainer

Pod Lifecycle



⇒ The phase of a pod is a simple high-level summary of where it is in its lifecycle.

Pending

⇒ The pod has been accepted by the k8s system, but it's not running.

⇒ One or more of the containers is still downloading.

⇒ if the pod cannot be scheduled because of resource constraints.

Running

⇒ The pod has been bound to a node.

⇒ All of the containers have been created.

⇒ At least one container is still running or is in the process of starting or restarting.

Succeeded

⇒ All containers in the pod have terminated in success, and will not be restarted.

Failed

⇒ All containers in the pod have terminated and at least one container has terminated in failure.

⇒ The Container either exited with non-zero status or was terminated by the System.

Completed

- ⇒ The pod has run to completion as there's nothing to keep it running
e.g - completed job.

Unknown

- ⇒ State of the pod could not be obtained.

- ⇒ Typically due to an error in network or communicating with the host of the pod.

Pod Conditions

A pod has a `podStatus`, which has an array of `podConditions` through which the pod has or has not passed.

- ⇒ Using `"kubectl describe pod PodName"` you can get condition of a pod.

These are the possible types:

podScheduled :— The pod has been scheduled to a node.

Ready :— The pod is able to serve request and will be added to the load balancing pool of all matching service.

Initialized :- All init Containers have started successfully

Unscheduled :- The scheduler cannot schedule the pod right now. for e.g due to lacking of Resource or other constraints.

Container Ready :- All containers in the pod are ready.